



ANSIBLE BEST PRACTICE



ANDREAS NEEB
SOLUTION ARCHITECT
RED HAT

aneeb@redhat.com

<http://github.com/andyneeb/presentations>

For God sakes...aren't best practices in 2019 dead?

Successful Automation





I ♥
ANSIBLE


```
cat baby.yml
---
- name: baby
  hosts: parental_units
  roles:
    - eat
    - sleep
    - poop
    - love
```

A

Even I can run
a playbook

A

A Powerful Team



**An enterprise-wide
automation strategy
must benefit
individuals first.**

Starting with the BIG picture is not the best path to enlightenment

Start the revolution from your desk

Solving smaller problems in repeatable fashion is easier to unify

Look for quick wins, current gaps

Make easy but noticeable progress

Map out orchestration, workflows etc

Part 1: Principles

#1: COMPLEXITY KILLS PRODUCTIVITY

That's not just a marketing slogan. We really mean it and believe that. We strive to reduce complexity in how we've designed Ansible tools and encourage you to do the same. **Strive for simplification in what you automate.**

#2: OPTIMIZE FOR READABILITY

If done properly, it can be the documentation of your workflow automation.

#3: THINK DECLARATIVELY

Ansible is a desired state engine by design. If you're trying to "write code" in your plays and roles, you're setting yourself up for failure. Our YAML-based playbooks were never meant to be for programming.


Part 2: Practices

KISS

- Keep plays and playbooks focused. Multiple simple ones are better than having a huge single playbook full of conditionals
- Once a playbook gets long or you're repeating tasks, use roles
- Follow Linux principle of do one thing, and one thing well

Step 1: Version control your Ansible content

- Start as simple as possible and iterate
 - Start with a basic playbook and static inventory
 - Refactor and modularize later
- Start with one Git repository - but when it grows, use multiple!




#2: Treat your Ansible content like code

ANSIBLE

Example: GitHub workflow

1. **Does not** require GitHub, the workflow model is just called that
2. **A** very simple workflow
3. **Master** branch is always possible to release
4. **Branches** are where you develop and test new features and bugfixes.
5. **Yes**, I wrote test. If you do not test your Ansible code you cannot keep the master branch releasable and this all fails.



Step 2 - N: CI / CD

- Verify correct syntax (`--syntax-check`)
- Verify style for bad practices (`ansible-lint`)
- Run your playbook or role and ensure it completes without failures, run again, check idempotency
- Test, test, test (automated!)

Give inventory nodes human-meaningful

EXHIBIT A

```
10.1.2.75  
10.1.5.45  
10.1.4.5  
10.1.0.40
```

```
w14301.example.com  
w17802.example.com  
w19203.example.com  
w19304.example.com
```



EXHIBIT B

```
db1  ansible_host=10.1.2.75  
db2  ansible_host=10.1.5.45  
db3  ansible_host=10.1.4.5  
db4  ansible_host=10.1.0.40
```

```
web1  ansible_host=w14301.example.com  
web2  ansible_host=w17802.example.com  
web3  ansible_host=w19203.example.com  
web4  ansible_host=w19203.example.com
```

Group hosts for easier inventory selection and less conditional tasks -- the more groups the better.

WHAT

```
[db]  
db[1:4]
```

```
[web]  
web[1:4]
```

```
db1 = db, east, dev
```

WHERE

```
[east]  
db1  
web1  
db3  
web3
```

```
[west]  
db2  
web2  
db4  
web4
```

WHEN


```
[dev]  
db1  
web1
```

```
[test]  
db3  
web3
```

```
[prod]  
db2  
web2  
db4  
web4
```

Use dynamic sources where possible. Either a single source of truth or let Ansible unify them.

- Stay in sync automatically
- Reduce human error
- No lag when changes occur
- Let others manage the inventory



No!

- `name: install telegraf`
 `yum: name=telegraf-{{ telegraf_version }} state=present update_cache=yes disabl`
 `notify: restart telegraf`
- `name: configure telegraf`
 `template: src=telegraf.conf.j2 dest=/etc/telegraf/telegraf.conf`
- `name: start telegraf`
 `service: name=telegraf state=started enabled=yes`

Yes!

```
- name: install telegraf
  yum:
    name: telegraf-{{ telegraf_version }}
    state: present
    update_cache: yes
    disable_gpg_check: yes
    enablerepo: telegraf
  notify: restart telegraf

- name: configure telegraf
  template:
    src: telegraf.conf.j2
    dest: /etc/telegraf/telegraf.conf
  notify: restart telegraf
```


Don't just start services -- use smoke tests

```
- name: check for proper response
  uri:
    url: http://localhost/myapp
    return_content: yes
  register: result
  until: '"Hello World" in result.content'
  retries: 10
  delay: 1
```

Separate provisioning from deployment and configuration tasks

```
acme_corp/  
├── configure.yml  
├── provision.yml  
└── site.yml
```

```
$ cat site.yml  
---  
- import_playbook: provision.yml  
- import_playbook: configure.yml
```

$$f(f(x)) = f(x)$$

- Use the `run` `command` modules like *shell* and *command* as a last resort
- The `command` module is generally safer
- The `shell` module should only be used for I/O redirect

Still using `command` a lot? Develop your own modules

The world is flat - Proper variable naming can make plays more readable and avoid variable name conflicts

- Use descriptive, unique human-meaningful variable names
- Prefix role variables with its “owner” such as a role name or package

```
apache_max_keepalive: 25
apache_port: 80
tomcat_port: 8080
```

- Do not use every possibility to store variables - settle to a defined scheme and as few places as possible
- Document in /defaults

No!

```
- hosts: web
  tasks:
    - yum:
        name: httpd
        state: latest

    - service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [web]
*****

TASK [setup]
*****
ok: [web1]

TASK [yum]
*****
ok: [web1]

TASK [service]
*****
ok: [web1]
```

Yes!

```
- hosts: web
  name: install and start apache
  tasks:
    - name: install apache packages
      yum:
        name: httpd
        state: latest

    - name: start apache service
      service:
        name: httpd
        state: started
        enabled: yes
```

```
PLAY [install and start apache]
*****

TASK [setup]
*****
ok: [web1]

TASK [install apache packages]
*****
ok: [web1]

TASK [start apache service]
*****
ok: [web1]
```


Careful when mixing manual and automated configuration
(Or even different automation frameworks...)

- Label template output files as being generated by Ansible

```
{{ ansible_managed | comment }}
```

```
#  
# Ansible managed  
#  
search example.com  
nameserver 192.168.122.1
```

Keep in mind

- Like playbooks -- keep roles purpose and function focused
- Use a `roles/` subdirectory for roles developed for organizational clarity in a single project
- Follow the Ansible Galaxy pattern for roles that are to be shared beyond a single project
- Limit role dependencies

Tricks and tips

- Use `ansible-galaxy init` to start your roles...
- ...then remove unneeded directories and stub files
- Use `ansible-galaxy` to install your roles -- even private ones
- Use a roles files (i.e. `requirements.yml`) to manifest any external roles your project is using

Thank you

Complexity kills productivity

Optimize for readability

Think declaratively